

TOKEN SPEND

Two ways to cut tokens

Only one of them moves the bill

TWO SCHOOLS

- Caveman cuts the output
- Agent answers telegraphic
- RTK cuts the input
- Tools return a digest

THE CLAIM SHRANK

- 75% — author's own test
- 65% — on 10 chat prompts
- 8.5% — on 86 real tasks
- plus 1.5k input tokens per turn

INDEPENDENT TEST

8.5%

Caveman real saving

JETBRAINS · 86 AGENTIC TASKS

THE OTHER SIDE

87.7%

RTK on my harness

477M TOKENS · 114K COMMANDS

TAKEAWAY

**Cut what the
model reads**

github.com/JuliusBrussee/caveman