

HINDSIGHT

Agent memory that changes its mind

Not chat chunks, but beliefs with a
confidence score

THE PROBLEM IT SOLVES

- every project starts your agent from zero
- it recalls text, never conclusions
- burned once, forgotten by morning
- memory grows, behaviour stays the same

HOW IT WORKS

- each new fact is checked against beliefs
- reinforce / weaken / contradict
- confidence moves one step, never overwritten
- a contradiction costs a double step

ON THE BENCHMARK

94.6%

LongMemEval, top of the board

Mem0 claims 94.4% — that gap is noise

WHAT IT TAKES TO RUN

- no database needed, Postgres is bundled
- embeddings run locally, no API key
- LLM can ride your existing subscription
- one MCP endpoint per memory bank

THE HONEST TRADE-OFFS

Cloud retain costs \$10 per million tokens — more than the model itself. A \$150-per-hour bug was fixed days ago.

MORE

**Worth it not for the score,
but for memory that
knows how to doubt**

t.me/aidacode